

# Changing Data in BigQuery

04

# Topics

- |    |                                       |
|----|---------------------------------------|
| 01 | Managing Change in Data Warehouses    |
| 02 | Handling Slowly Changing Dimensions   |
| 03 | DML Statements in BigQuery            |
| 04 | DML: Best Practices and Common Issues |



A blue square containing the white text '01'. To the right of the square is a yellow triangle pointing downwards.

01

# Managing Change in Data Warehouses

# BigQuery manages service availability, but you control the duration and timeliness of your datasets

- All table modifications are ACID-compliant.
- Timeliness might be affected by:
  - Periodic loads versus streaming ingest
  - Priority of load jobs versus analytics jobs

# No need to delete older data

In BigQuery, you can afford to retain older data:

- Use sliding windows based on partition field.
- Benefit from BigQuery long-term storage pricing.


If deleting older data is necessary, set up automatic table/partition expiry.

# Make changes without affecting users

- Data warehouses are constantly evolving as users want new and different functionality and data.
- As a result, you must be able to change both schemas and data without affecting users.

# Schema changes are usually scheduled as version upgrades

Design, develop, and test the upgrade in parallel while the previous version of the data warehouse is serving analysis workloads.



# What about data changes?

Typically *facts*  
don't change,  
but  
*dimensional*  
attributes  
might.



# Fact tables versus dimension tables

## Fact table

holds measurements, metrics, or facts about a business process.

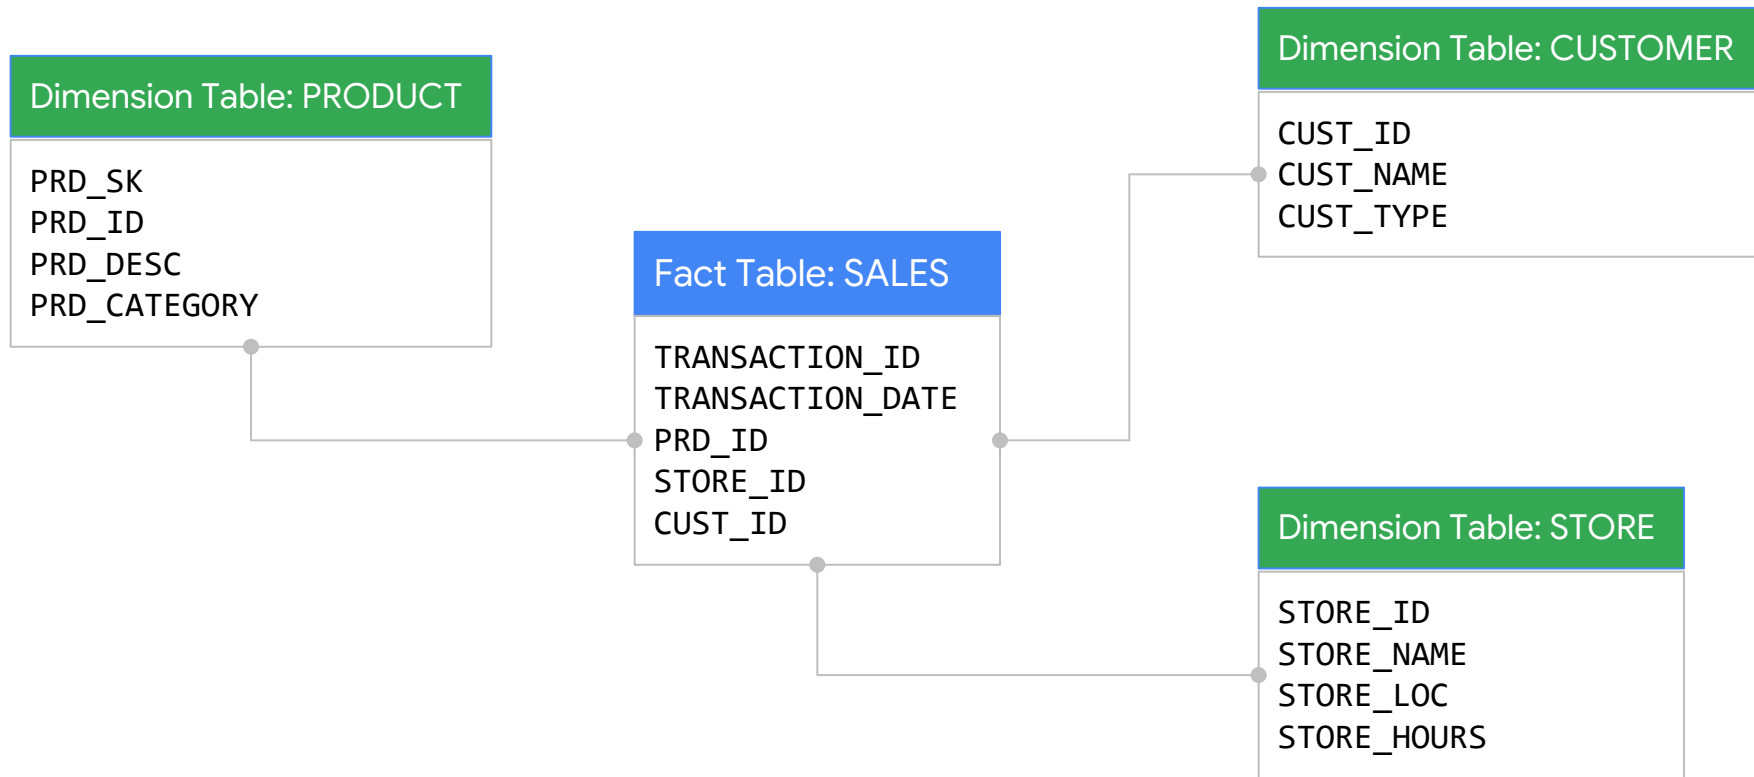
- Examples include sales records.
- Has two types of columns: those that contain facts and those that are a foreign key to a dimension table.

## Dimension table

contains descriptive attributes to be used as query constraining.

- Examples include customer, employee, product data.
- Relatively static data **that can change slowly** but unpredictably, rather than according to a regular schedule.

# Star schema example with fact and dimension tables



# Slowly changing dimensions

In data warehousing,  
**Slowly Changing Dimensions (SCD)**  
is an important concept  
that is used to enable the  
*historic aspect of data*  
in an analytical system.



# Handling Slowly Changing Dimensions

# SCD Type 1: Overwrite attribute value

Before:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY
123	ABC	awesome moisturizer cream - 100 oz	health and beauty

After:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY
123	ABC	awesome moisturizer cream - 100 oz	<del>health and beauty</del> cosmetics

# SCD Type 2: Change attribute value and maintain history

Before:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY	START_DATE	END_DATE
123	ABC	ace moisturizer cream - 100 oz	health and beauty	31-Jan-2009	NULL

After:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY	START_DATE	END_DATE
123	ABC	ace moisturizer cream - 100 oz	health and beauty	31-Jan-2009	18-JUL-2017
124	ABC	ace moisturizer cream - 100 oz	cosmetics	19-JUL-2017	NULL

# Create view to use in analytics queries

```
CREATE VIEW products_current as (  
    SELECT PRD_SK, PRD_ID, PRD_DESC, PRD_CATEGORY, PRD_START_DATE  
    FROM dimension_table  
    WHERE END_DATE IS NULL  
);
```

In a denormalized schema, no changes may be needed to previous fact table rows

TRANSACTION_DATE	PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY	UNITS	AMOUNT
18-JUL-2017	123	ABC	ace moisturizer cream - 100 oz	health and beauty	2	25.16
19-JUL-2017	124	ABC	ace moisturizer cream - 100 oz	cosmetics	1	13.50



# SCD Type 2: Other options

Using version number to maintain history:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY	VERSION
123	ABC	ace moisturizer cream - 100 oz	health and beauty	0
124	ABC	ace moisturizer cream - 100 oz	cosmetics	1

Using effective date and current flag:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY	EFFECTIVE_DATE	CURRENT_FLAG
123	ABC	ace moisturizer cream - 100 oz	health and beauty	31-Jan-2009	N
124	ABC	ace moisturizer cream - 100 oz	cosmetics	19-JUL-2017	Y

# SCD Type 3: Maintain history by adding columns

Base table:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY		
123	ABC	ace moisturizer cream - 100 oz	CATEGORY_NAME	START_DATE	END_DATE
			health and beauty	31-Jan-2009	18-JUL-2017
			cosmetics	18-JUL-2017	NULL

# Create view to use in analytics queries

```
CREATE VIEW products_current as (  
  SELECT PRD_SK, PRD_ID, PRD_DESC,  
         PRD_CATEGORY.ordinal[array_length(PRD_CATEGORY)] as PRD_CAT  
  FROM dimension_table  
);
```

View:

PRD_SK	PRD_ID	PRD_DESC	PRD_CAT
124	ABC	ace moisturizer cream - 100 oz	cosmetics


# For SCD, there's no one-size-fits-all solution

Changes may be handled differently due to the performance implications of DML:


- In a denormalized table:
  - View switching
  - In-place partition loading
- Isolate changes by using normalized data schema and isolate change in small dimension tables.
  - Update data masking.

# Option 1: View switching


## 1. Allocate new table




## 2. Load new data



## 3. Redefine view



## 4. Deallocate old table



# View switching considerations

- Zero down-time
- DML on main view might be problematic
- Apply view-switching to custom partitions

## Option 2: In-place partition loading

To replace data in a target partition with data from a query of another table:

```
bq query --use_legacy_sql=false --replace \  
  --destination_table 'flight_data.fact_flights_part$20140910' \  
  'select * REPLACE(...) from `ods.load_flights_20140910`'
```

To replace data in a target partition by loading from Cloud Storage:

```
bq load --replace \  
  --source_format=NEWLINE_DELIMITED_JSON \  
  'flight_data.fact_flights_part$20140910' \  
  gs://{bucket}/load_flights_20140910.json
```

## Option 3: Update data masking

For tables with data that can change frequently, even within the course of a day, you can implement a conditional join through a view.

```
SELECT f.order_id as order_id, f.customer_id as customer_id,  
       IFNULL(u.customer_first_name, f.customer_first_name) as customer_first_name,  
       IFNULL(u.customer_last_name, f.customer_last_name) as customer_last_name  
FROM fact_table f  
LEFT OUTER JOIN pending_customer_updates u  
ON f.customer_id = u.customer_id
```



03



## DML Statements in BigQuery

# Overwrite with UPDATE DML statement for SCD Type 1

```
UPDATE dimension_table  
SET PRD_CATEGORY="cosmetics"  
WHERE PRD_SK="123"
```

Dimension table:

PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY
123	ABC	awesome moisturizer cream - 100 oz	<del>health and beauty</del> cosmetics

# Maintain row history using INSERT for SCD Type 2

Insert new record into dimension table :

```
INSERT dimension_table  
(PRD_SK, PRD_ID, PRD_DESC, PRD_CATEGORY, VERSION)  
VALUES (124,'ABC','ace moisture cream - 100 oz','cosmetics',1)
```

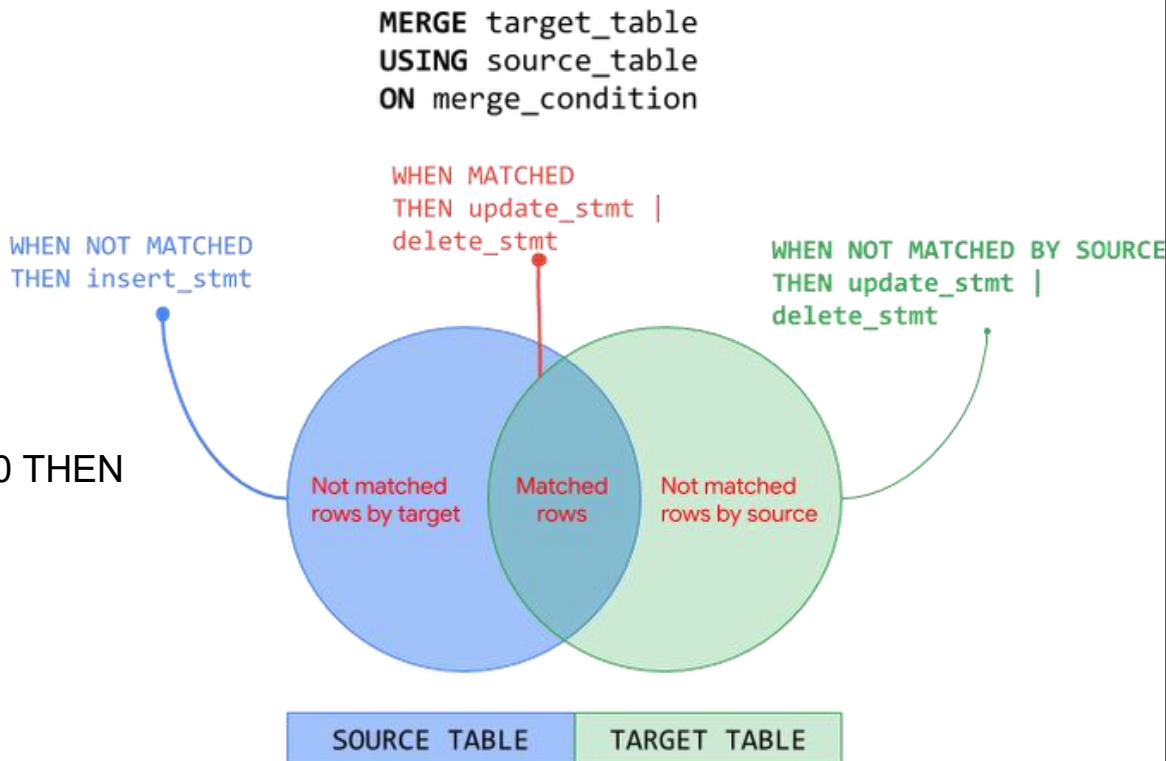
PRD_SK	PRD_ID	PRD_DESC	PRD_CATEGORY	VERSION
123	ABC	ace moisturizer cream - 100 oz	health and beauty	0
124	ABC	ace moisturizer cream - 100 oz	cosmetics	1

# UPDATE = DELETE + INSERT

- UPDATE is implemented as **DELETE old row** + **INSERT updated row**.
- For SCD Type 2, where you change the attribute value to maintain history, you do an UPDATE and an INSERT.
- For example, you will *update* the end\_date or current\_flag attribute of the existing row, and *insert* a new row with a null end\_date or Y flag.

# MERGE is INSERT, UPDATE, DELETE over a single table in one statement

```
MERGE target_table T
USING source_table S
ON T.field2 = S.field2
WHEN MATCHED AND T.field2 > 100 THEN
  DELETE
WHEN NOT MATCHED THEN
  INSERT(field2) VALUES(field2)
```



# Update dimension table using MERGE UPDATE for SCD Type 1

```
MERGE dimension_table as MAIN using  
temporary_table as TEMP  
on MAIN.PRD_SK = TEMP.PRD_SK  
when matched then  
UPDATE SET  
MAIN.PRD_CATEGORY = TEMP.PRD_CATEGORY  
when not matched then  
INSERT VALUES (TEMP.PRD_SK, TEMP. PRD_ID, TEMP. PRD_SK, TEMP.  
PRD_CATEGORY)
```

# DML concurrency

- During any 24-hour period, the first 1500 statements that INSERT into a table run concurrently.
- The UPDATE, DELETE, and MERGE DML statements are called *mutating DML statements*.
  - BigQuery runs up to 2 of them concurrently, after which up to 20 are queued as PENDING.

# DML conflicts

- Mutating DML statements that run concurrently on a table cause DML statement conflicts when the statements try to mutate the same partition.
- The statements succeed only if they don't modify the same partition. BigQuery tries to rerun failed statements up to three times.





## DML: Best Practices and Common Issues

# DML best practices

- Use partitioned tables if updates or deletions generally happen on older data or on data in a date-localized manner.
- Use updates over clustered tables where there is clustering locality on the modified rows; they will generally perform better.
- Group DML operations together into larger ones to amortize the cost of running smaller operations.
- Avoid partitioning tables if the amount of data in each partition is small and each update modifies a large fraction of partitions in the table.

# DML considerations

- Maximum of 2 concurrent non-insert DML jobs per table. This helps prevent conflicts.
  - After which, BigQuery queues up to 20 non-insert DML jobs per table.
  - Additional statements past the maximum queue length for each table fail.
- Metadata update rate limit is applied at job insert time.
- Commit latency time after the job is done can be very long sometimes.
  - This gets worse with a large number of partitions.
- Single row inserts or single row updates are generally an **undesirable pattern**.

# DML common performance bottlenecks

- Updating rows with no locality causes rewrite of entire table, even with partitions.
  - For example, when you update few rows from all files.
  - Consider using clustering.
- Commit time sometimes runs longer than actual execution time.
  - “Byte Counting Query” is required for Billing Quota Check
- Narrow mutations (or deletions) with locality not as fast as they should be because the underlying files are large.

# Questions?

