

Query Processing and Optimization

06

Topics


01 Simple Query Execution

02 SELECTs and Aggregation

03 JOINs and Skewed JOINs

04 Filtering and Ordering

05 Best Practices for Functions




01



Simple Query Execution

Simple query execution



```
SELECT COUNT(*)  
FROM `new_york.citibike_trips`  
WHERE start_station_name LIKE  
"%Broadway%"
```


Stage 2: Sum (1 worker)

Stage 1: Filter, count (220 inputs)

Demo

Simple query execution

Query plan, stages, and metrics



Interpreting the query plan

- **Is there a significant difference between avg and max time?**
 - There is probably a data skew; use APPROX_TOP_COUNT to check.
 - Filter early to work around.
- **Most time is spent reading from intermediate stages.**
 - Consider filtering earlier in the query.
- **Most time is spent on CPU tasks.**
 - Consider approximate functions, inspect UDF usage, filter earlier.

How do you optimize queries?

Less work → Faster query

What is **work** for a query?

- **I/O** — How many bytes did you read?
- **Shuffle** — How many bytes did you pass to the next stage?
- **Grouping** — How many bytes do you pass to each group?
- **Materialization** — How many bytes did you write?
- **CPU work** — User-defined functions (UDFs), functions



02

SELECTs and Aggregations

SELECTs

Optimization: Necessary columns only

Original code

```
select  
  *  
from  
`dataset.table`
```

Optimized

```
select  
  * EXCEPT (dim1, dim2)  
from  
`dataset.table`
```

Reasoning

Only select the columns necessary, especially in inner queries.
SELECT * is cost-inefficient and may also hurt performance.

If the number of columns to return is large, consider using
SELECT * EXCEPT to exclude unneeded columns.

In some use cases, **SELECT * EXCEPT** may be necessary.

Recap: Auto-pruning with partitioning and clustering

Partitioned table

Table info

Table ID	fh-bigquery:wikipedia_v2.pageviews_2017
Table size	2.2 TB
Long-term storage size	2.2 TB
Number of rows	54,489,325,868
Created	Feb 27, 2018, 1:54:41 AM
Table expiration	Never
Last modified	Feb 27, 2018, 4:47:50 AM
Data location	US
Table type	Partitioned
Partitioned by	Day
Partitioned on field	datehour
Partition filter	Required

Partitioned and clustered table

Table info

Table ID	fh-bigquery:wikipedia_v3.pageviews_2017
Table size	2.2 TB
Long-term storage size	2.2 TB
Number of rows	54,489,325,868
Created	Aug 1, 2018, 1:24:57 AM
Table expiration	Never
Last modified	Aug 2, 2018, 8:50:32 PM
Data location	US
Table type	Partitioned
Partitioned by	Day
Partitioned on field	datehour
Partition filter	Required
Clustered by	wiki, title

Recap: Auto-pruning with partitioning and clustering

Partitioned table by datehour

```
SELECT *  
FROM `fh-bigquery.wikipedia_v2.pageviews_2017`  
WHERE DATE(datehour) BETWEEN '2017-06-01'  
AND '2017-06-30'  
LIMIT 1
```

1.7 sec elapsed, 180 GB processed

Partitioned table by datehour

```
Clustered table by wiki, title  
SELECT *  
FROM `fh-bigquery.wikipedia_v3.pageviews_2017`  
WHERE DATE(datehour) BETWEEN '2017-06-01'  
AND '2017-06-30'  
LIMIT 1
```

1.8 sec elapsed, 16 MB processed

Nest repeated data

- Customers often default to “flat” denormalization even if it is not the most beneficial.
 - Requires a GROUP BY to analyze data.
- Example: Orders table with a row for each line item
 - {order_id1, item_id1}, {order_id1, item_id2}, ...
- If you model one order per row and nest line items in a nested field, GROUP BY is no longer required.
 - {order_id1, [{item_id1}, {item_id2}] }

Query cache

Hash of

- Data modification times
- Tables used
- Query string

Cache skipped if

- Referenced tables or views have changed
- Non-deterministic function used (e.g., NOW())
- Permanent result table requested
- Source tables have streaming buffers

Hash becomes output table name

- Cache is per-user

Query settings

Query engine

BigQuery engine

Cloud Dataflow engine

Deploy your data processing pipelines on the Cloud Dataflow service.

Destination

Set a destination table for query results

Project name

danny-bq

Dataset name

big_query_testing

Table name

Letters, numbers, and underscores allowed

Advanced options ▾

Resource management

Job priority

Interactive

Batch

Cache preference

Use cached results

Aggregation

Optimization: Late aggregation

Original code

```
select
    t1.dim1,
    sum(t1.m1),
    sum(t2.m2)
from (select
        dim1,
        sum(metric1) m1
    from `dataset.table1` group by 1) t1
join (select
        dim1,
        sum(metric2) m2
    from `dataset.table2` group by 1) t2
on t1.dim1 = t2.dim1
group by 1;
```

Optimized

```
select
    t1.dim1,
    sum(t1.m1),
    sum(t2.m2)
from (select
        dim1,
        metric1 m1
    from `dataset.table1`) t1
join (select
        dim1,
        metric2 m2
    from `dataset.table2`) t2
on t1.dim1 = t2.dim1
group by 1;
```

Reasoning


Aggregate as late and as seldom as possible because **aggregation is very costly**.

BUT if a table can be reduced drastically by aggregation in preparation for being joined, then aggregate it early.

Caution: With JOINS, this only works if the two tables are already aggregated to the same level (i.e., if there is only one row for every join key value).


Shuffle aggregation execution

Simple select query



VS

Aggregation query



```
SELECT language, MAX/views  
as views  
FROM  
`wikipedia_benchmark.Wiki1B`  
WHERE title LIKE "G%o%"  
GROUP BY 1 ORDER BY 2 DESC  
LIMIT 100
```

Stage 3: SORT,
LIMIT (1 worker)

Stage 2: GROUP BY,
SORT, LIMIT (289 workers)


Shuffle

Stage 1: Partial
GROUP BY (40,859 sinks)

Demo


Aggregation with shuffle query
plan

Query plan, stages, and metrics




Aggregation with high cardinality

Without shuffle (non-BQ)




With Shuffle (BQ)



```
SELECT foo,
       COUNT(*) as cnt
  FROM `...`
 GROUP BY 1
 ORDER BY 2 DESC
 LIMIT 2
```

- Can't discard "B" or "C" until after all previous stages are complete.
- Shuffle puts like values in the same node.
- Scalable because you never have to return more than N from each node in middle tier.


Repartitioning



Demo

Repartitioning

Query plan, stages, and metrics



03



JOINs and Skewed JOINs

JOINs

Optimization: JOIN pattern

Original code

```
select
    t1.dim1,
    sum(t1.metric1),
    sum(t2.metric2)
from
    small_table t1
join
    large_table t2
on
    t1.dim1 = t2.dim1
where t1.dim1 = 'abc'
group by 1;
```

Optimized

```
select
    t1.dim1,
    sum(t1.metric1),
    sum(t2.metric2)
from
    large_table t2
join
    small_table t1
on
    t1.dim1 = t2.dim1
where t1.dim1 = 'abc'
group by 1;
```

Reasoning

When you create a query by using a JOIN, consider the **order** in which you are merging the data. The standard SQL query optimizer can determine which table should be on which side of the join, but we still recommend that you order your joined tables appropriately.

The best practice is to manually **place the largest table first**, followed by the smallest, and then by decreasing size. Only under specific table conditions does BigQuery automatically reorder/optimize based on table size.

Optimization: Filter before JOINs

Original code

```
select
    t1.dim1,
    sum(t1.metric1),
    sum(t2.metric3)
from
    `dataset.table1` t1
left join
    `dataset.table2` t2
on
    t1.dim1 = t2.dim1
where t2.dim2 = 'abc'
group by 1;
```

Optimized


```
select
    t1.dim1,
    sum(t1.metric1),
    sum(t2.metric3)
from
    `dataset.table1` t1
left join
    `dataset.table2` t2
on
    t1.dim1 = t2.dim1
where t1.dim2 = 'abc' AND t2.dim2 = 'abc'
group by 1;
```

Reasoning


WHERE clauses should be executed **as soon as possible, especially within joins**, so the tables to be joined are as small as possible.

WHERE clauses may not always be necessary, because standard SQL will do its best to push down filters. Review the explanation plan to see whether filtering is happening as early as possible, and either fix the condition or use a subquery to filter in advance.

Small JOIN (broadcast)



Large JOIN (shuffle)



Hash join


Independent shuffles

```
SELECT c.author.name a, c2.a m  
FROM github_repos.commits c  
JOIN (SELECT committer.name a, commit  
      FROM github_repos.commits) c2  
ON c.commit = c2.commit  
LIMIT 1000
```

Demo

Broadcast JOIN and Shuffle
JOIN

Query plan, stages, and metrics



JOIN optimization with clustered tables

Subquery titles (“Animated_Google”, “A_Google_A_Day”, etc.)
are used to filter table `pageviews_2017` BEFORE joining.

```

SELECT
  tbl2017.*
FROM
  `wikipedia_v3.pageviews_2017` `tbl2017`
JOIN(SELECT * FROM
  `wikipedia_vt.just_latest_rows` `
  WHERE
    REGEXP_CONTAINS(title, "Google"))
USING(title)
WHERE
  DATE(tbl2017.datehour)
  BETWEEN '2017-06-01' AND '2017-06-30'

```

Query complete (0.8 sec elapsed, 21.1 MB processed)

Job information	Results	JSON	Execution details	
Row	datehour	wiki	title	views
1	2018-10-22 16:00:00 UTC	en	Animated_Google	1
2	2018-10-23 14:00:00 UTC	en	Alphabet_(Google)	1
3	2018-10-24 23:00:00 UTC	en	Actions_on_Google	1
4	2018-10-24 23:00:00 UTC	en	Authors_Guild,_Inc._v._Google,_Inc.	2
5	2018-10-22 15:00:00 UTC	en	Acquisitions_by_Google	1
6	2018-10-23 02:00:00 UTC	en	Adblock_Plus_for_Google_Chrome	1
7	2018-10-23 15:00:00 UTC	en	Accomplished_Googlebombs	1
8	2018-10-24 16:00:00 UTC	en	Are_You_SmartEnough_to_Work_at_Google?	1
9	2018-10-24 21:00:00 UTC	en	Authors_Guild_v._Google	1
10	2018-10-24 21:00:00 UTC	en	A_Google_A_Day	1


Clustered table JOIN optimization requirements

```
SELECT  
  tb12017.*  
FROM  
  `wikipedia_v3.pageviews_2017`  tbl2017  
JOIN(SELECT * FROM  
    `wikipedia_vt.just_latest_rows`  
    WHERE  
      REGEXP_CONTAINS(title, "Google"))  
  USING(title)  
WHERE  
  DATE(tbl2017.datehour)  
  BETWEEN '2017-06-01' AND '2017-06-30'
```


Left table `pageviews_2017` must be clustered.

Subquery result size must qualify for **broadcast join** (e.g., `JOIN EACH WITH ALL`).

Clustered table JOIN optimization in query plan




Clustered table JOIN optimization in query plan



Partitioned but **not clustered** table results in joining **more data** and consuming **more slots**.

Partitioned **and clustered** table results in joining **less data** and consuming **fewer slots**.




Join explosions

- Are caused by JOIN with non-unique key on both sides.
- SQL relational algebra gives cartesian product of rows that have the same join key.
 - Worst case: The number of output rows is the number of rows in the left table multiplied by the number of rows in the right table.
 - In extreme cases, the query will not finish.
- If the job finishes, the query explanation will show output rows versus input rows.
- Confirm the diagnosis by modifying the query to print the number of rows on each side of the JOIN grouped by the JOIN key.
- A workaround is to use GROUP BY to pre-aggregate.

Skewed JOINs

Skewed JOIN




One worker gets too much data

Independent shuffles


Demo

Skewed JOIN

Query plan, stages, and metrics



Skewed JOIN query plan



Skewed JOINS

- **BigQuery shuffles data on each side of the join:**
 - All data with the same join key goes to the same worker.
 - Data can overload the worker.
- **Typically result from data skew**
- **Workarounds:**
 - Pre-filter rows from query with the unbalanced key.
 - Potentially split into two queries.



Filtering and Ordering

WHERE clause: Expression order matters!

Original code

```
SELECT text  
FROM  
`stackoverflow.comments`  
WHERE  
text LIKE '%java%'  
AND user_display_name = 'anon'
```

Optimized

```
SELECT text  
FROM  
`stackoverflow.comments`  
WHERE  
user_display_name = 'anon'  
AND text LIKE '%java%'
```


The expression:
user_display_name = 'anon'
filters out much more data
than the expression:
text LIKE '%java%'

Reasoning

BigQuery assumes that the user has provided the best order of expressions in the WHERE clause and does not attempt to reorder expressions. Expressions in your WHERE clauses should be ordered with the most selective expression first.

The optimized example is faster because it doesn't execute the expensive LIKE expression on the entire column content, but instead only on the content from user 'anon'.

WHERE clause reordering: Proof in the query plan



```
WHERE
  text LIKE
  '%java%'
  AND
  user_display_name
=
  'anon'
```

```
WHERE
  user_display_name
=
  'anon'
  AND
  text LIKE
  '%java%'
```

Optimization: ORDER BY with LIMIT

Original code

```
select  
    t.dim1,  
    t.dim2,  
    t.metric1  
from  
    `dataset.table` t  
order by t.metric1 desc
```

Optimized


```
select  
    t.dim1,  
    t.dim2,  
    t.metric1  
from  
    `dataset.table` t  
order by t.metric1 desc  
limit 1000
```

Reasoning

Writing results for a query with an **ORDER BY** clause can result in **Resources Exceeded** errors. Because the final sorting must be done on a single slot, if you are attempting to order a very large result set, the final sorting can overwhelm the slot that is processing the data.

If you are sorting a very large number of values, use a **LIMIT** clause.


Large ORDER BYs



```
SELECT foo  
FROM table  
ORDER BY foo
```

Master node needs to
sort and store all
values


ORDER BY and LIMIT



```
SELECT foo  
FROM table  
ORDER BY foo  
LIMIT 3
```

Can drop values
over the limit at
each node

Overloaded ORDER BY query plan



```

SELECT
    version
  FROM
    bitcoin_blockchain.
    transactions
  ORDER BY
    version DESC
  
```

ORDER BY with LIMIT query plan



```

SELECT
    version
  FROM
    bitcoin_blockchain.
    transactions
  ORDER BY
    version DESC
  LIMIT 1000
  
```

05



Best Practices for Functions

Optimization: String comparison

Original Code

```
select  
    dim1  
from  
    `dataset.table`  
where  
    regexp_contains(dim1, '.*test.*')
```

Optimized

```
select  
    dim1  
from  
    `dataset.table`  
where  
    dim1 like '%test%'
```

Reasoning

REGEXP_CONTAINS > LIKE
where > means more functionality,
but also slower execution time.
Prefer LIKE when the full power of
regex is not needed (e.g., wildcard
matching).

Optimization: Approximate functions

Original code

```
select
  dim1,
  count(distinct dim2)
from
  `dataset.table`
group by 1;
```

Optimized

```
select
  dim1,
  approx_count_distinct(dim2)
from
  `dataset.table`
group by 1;
```

Reasoning

If the SQL aggregation function you're using has an equivalent approximation function, the approximation function will yield faster query performance.

Approximate functions produce a result which is generally **within 1%** of the exact number.

Optimization: SQL UDFs → JavaScript UDFs

Original code

```
create temporary function
    multiply(x INT64, y INT64)
returns INT64
language js
as """
    return x * y;
""";
```



```
select multiply(2, 2) as result;
```

Optimized

```
create temporary function
    multiply(x INT64, y INT64)
as
    (x * y);

select multiply(2, 2) as result;
```

Reasoning

JavaScript UDFs substantially reduce performance because they have to spin up a V8 subprocess evaluate.

Prefer SQL UDFs where possible.

Optimization: Persistent UDFs

The diagram illustrates the process of creating a persistent UDF and deploying it to a routine. On the left, a code editor shows the creation of a function:

```
CREATE OR REPLACE FUNCTION
your_dataset.addFourAndDivide(x INT64, y INT64) AS (
    (x + 4) / y
);
```

A blue arrow points from this code to the right side, where the resulting persistent function is shown in the Google Cloud console:

qwiklabs-gcp-00-d2dc922559ba

- External connections
- dts1
- Routines (1)
 - AddFourAndDivide

The "AddFourAndDivide" routine is highlighted with a blue arrow. Below this, the Google Cloud console interface shows the persistent function details:

AddFourAndDivide

Persistent function info

Persistent function ID	qwiklabs-gcp-00-d2dc922559ba.dts1.AddFourAndDivide
Created	Dec 11, 2022, 2:24:51 PM UTC-8
Last modified	Dec 11, 2022, 2:24:51 PM UTC-8
Language	SQL
Description	
Arguments	x INT64, y INT64
Return type	FLOAT64

Routine query

```
(x + 4) / y
```

INVOKED PERSISTENT FUNCTION EDIT PERSISTENT FUNCTION DELETE EDIT ROUTINE DETAILS

Google Cloud

Optimization: Persistent UDFs

Original code

```
CREATE TEMP FUNCTION addFourAndDivide(x  
INT64, y INT64) AS ((x + 4) / y);  
  
WITH numbers AS  
  (SELECT 1 as val  
  UNION ALL  
  SELECT 3 as val  
  UNION ALL  
  SELECT 4 as val  
  UNION ALL  
  SELECT 5 as val)  
SELECT val, addFourAndDivide(val, 2) AS  
result  
FROM numbers;
```

Optimized

```
# Replaced with persistent function  
# and invoked below  
  
WITH numbers AS  
  (SELECT 1 as val  
  UNION ALL  
  SELECT 3 as val  
  UNION ALL  
  SELECT 4 as val  
  UNION ALL  
  SELECT 5 as val)  
SELECT val,  
`your_project.your_dataset.addFourAndDivide`  
(val, 2) AS result  
FROM numbers;
```

Reasoning

Create persistent user-defined SQL and JavaScript functions in a centralized BigQuery dataset that can be invoked across queries and in logical views.

Create org-wide libraries of business logic within shared datasets.

Scripting and stored procedures

- Execute multiple statements in one request.
- Declare, assign, and use variables.
- Control execution with conditions and loops.
- **Caveat:** Statements are committed independently of each other.

```
1 DECLARE primes ARRAY<INT64> DEFAULT [2];
2 DECLARE n INT64 DEFAULT 3;
3 DECLARE max INT64 DEFAULT 30;
4 WHILE n <= max DO
5   BEGIN
6     DECLARE n_is_prime BOOL DEFAULT TRUE;
7
8     -- Test all prime numbers from 2 to SQRT(n), inclusive.
9     DECLARE i INT64 DEFAULT 0;
10    WHILE i < ARRAY_LENGTH(primes) DO
11      BEGIN
12        DECLARE prime INT64 DEFAULT primes[OFFSET(i)];
13        IF MOD(n, prime) = 0 THEN
14          -- Found a prime < n, which divides evenly into n, so n is not prime.
15          SET n_is_prime = FALSE;
16          BREAK;
17        END IF;
18        IF prime * prime >= n THEN
19          -- <primes> is kept sorted in increasing order, so once we find a
20          -- single value >= SQRT(n), we can stop.
21          BREAK;
22        END IF;
23      END;
24      SET i = i + 1;
25    END WHILE;
26
27    -- If n is prime, then add it to the list of known primes.
28    IF n_is_prime THEN
29      SET primes = ARRAY_CONCAT(primes, [n]);
30    END IF;
31  END;
32  SET n = n + 1;
33 END WHILE;
34
35 -- Display all the primes.
36 SELECT prime FROM UNNEST(primes) AS prime ORDER BY prime;
```

Questions?

Lab (75 min)

Optimizing Your BigQuery Queries for Performance

(choose your level)

1:15:00